

# SPL Technical Architecture v3.1

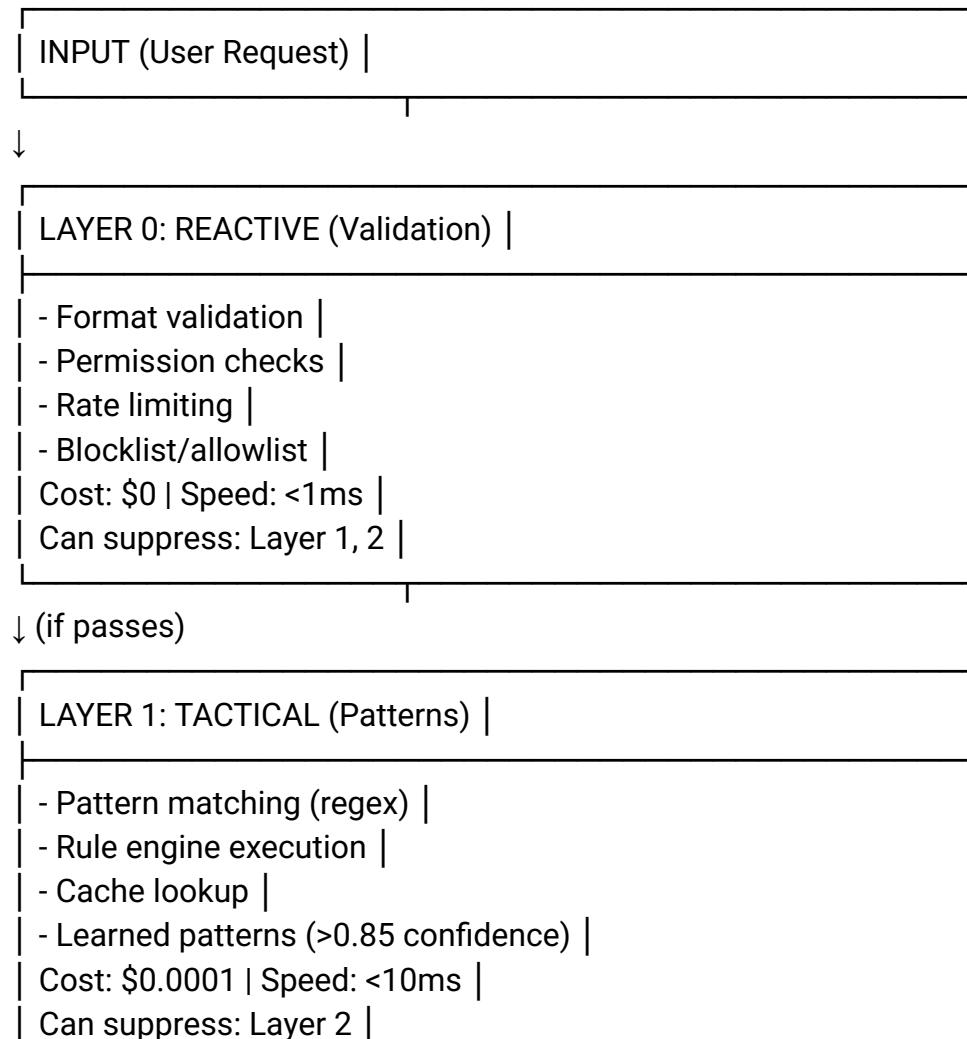
\*\*Implementation Guide for Three-Layer Hierarchical Agent Systems\*\*

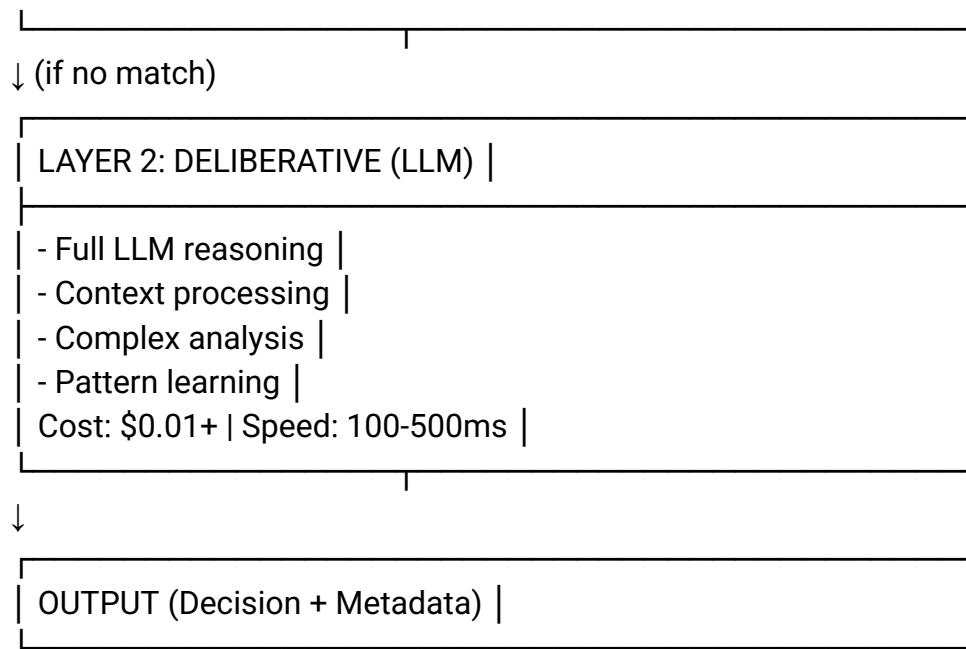
\*\*Author:\*\* Pamela Cuce  
\*\*Email:\*\* pamela@dasein.works  
\*\*Phone:\*\* 914-240-3564  
\*\*Version:\*\* 3.1 (Corrected)  
\*\*Date:\*\* December 3, 2025

---

## ## Architecture Overview

SPL implements a three-layer hierarchical decision system where each layer can suppress upper layers based on cost, confidence, and safety constraints.





## Layer 0: Reactive Layer

### ### Purpose

Fast, deterministic validation. Rejects invalid inputs before processing.

### ***Layer 0 Reactive Validation***

***Purpose: Zero-cost structural validation before expensive operations***

***This layer acts as first-pass filter to reject invalid requests immediately***

class Layer0Reactive:

====

Layer 0 handles all structural validation:

- Format checking (is input well-formed?)
- Permission verification (is user allowed?)
- Rate limiting (has user exceeded quotas?)
- Blocklist/allowlist enforcement (is user authorized?)

### ### Implementation

All operations here are O(1) or O(log n), guaranteed <1ms  
Cost is always \$0 - no API calls or expensive operations

.....

```
def __init__(self):
    # blocklist: set of user IDs that are blocked (fast O(1) lookup)
    self.blocklist = set()
    # allowlist: if non-empty, only these users are allowed
    self.allowlist = set()
    # rate_limits: dict tracking request counts per user per minute
    self.rate_limits = {}
```

```
def validate(self, request):
```

.....

Validate request structure

Args:

request: dict with 'user\_id' and 'content' keys

Returns:

tuple: (is\_valid: bool, error\_message: str or None)

Checks:

- Is request a dict? (not None, string, list, etc.)
- Does it have 'user\_id'? (required for authentication)
- Does it have 'content'? (required for processing)

.....

```
# Type check: request must be dict
```

```
if not isinstance(request, dict):
```

```
    return False, "Invalid request format"
```

```
# Check for required user_id field
```

```
if 'user_id' not in request:
```

```
    return False, "Missing user_id"
```

```
# Check for required content field
if 'content' not in request:
    return False, "Missing content"

# All checks passed
return True, None

def check_permissions(self, user_id):
    """
    Check if user has permission to use system

    Args:
        user_id: string identifier for user

    Returns:
        tuple: (has_permission: bool, error_message: str or None)

    Logic:
    1. If user in blocklist → REJECT (explicitly denied)
    2. If allowlist exists and user not in it → REJECT (only specific users)
    3. Otherwise → ALLOW (user authorized)
    """

    # Blocklist check: if user explicitly blocked, deny immediately
    if user_id in self.blocklist:
        return False, "User blocked"

    # Allowlist check: if allowlist configured and user not in it, deny
    if self.allowlist and user_id not in self.allowlist:
        return False, "User not allowed"

    # User passed all permission checks
    return True, None
```

```
def check_rate_limit(self, user_id):
    """
    Check if user has exceeded rate limit

    Args:
```

user\_id: string identifier for user

Returns:

tuple: (within\_limit: bool, error\_message: str or None)

Implementation:

- Track requests per user per minute
- Limit: 100 requests/minute (configurable)
- Time window: 60 seconds (1 minute)
- After 60s elapses, counter resets

.....

```
limit = 100 # Max 100 requests per minute per user
```

```
# Initialize tracking for new user
```

```
if user_id not in self.rate_limits:
```

```
    self.rate_limits[user_id] = {  
        'count': 0,  
        'timestamp': __import__('time').time()  
    }
```

```
# Get current tracking entry
```

```
entry = self.rate_limits[user_id]
```

```
# Calculate elapsed time since window started
```

```
elapsed = __import__('time').time() - entry['timestamp']
```

```
# If more than 60 seconds passed, reset window
```

```
if elapsed > 60:
```

```
    entry['count'] = 0  
    entry['timestamp'] = __import__('time').time()
```

```
# Check if user hit rate limit
```

```
if entry['count'] >= limit:
```

```
    return False, f"Rate limit exceeded ({limit}/min)"
```

```
# User within limit - increment counter and allow
```

```
entry['count'] += 1
```

```
return True, None
```

```
def execute(self, request):
```

.....

Execute all Layer 0 checks in sequence

Args:

request: dict with user\_id and content

Returns:

dict with keys:

- halt: bool (True = stop processing)
- reason: str error message if halted
- cost: float (always 0.0 for Layer 0)

Execution flow:

1. Validate structure
2. Check permissions
3. Check rate limit

If any check fails, return halt=True immediately

.....

# Check 1: Structure validation

valid, error = self.validate(request)

if not valid:

    return {'halt': True, 'reason': error, 'cost': 0.0}

# Check 2: Permission check

valid, error = self.check\_permissions(request['user\_id'])

if not valid:

    return {'halt': True, 'reason': error, 'cost': 0.0}

# Check 3: Rate limit check

valid, error = self.check\_rate\_limit(request['user\_id'])

if not valid:

    return {'halt': True, 'reason': error, 'cost': 0.0}

# All checks passed - safe to proceed to Layer 1

return {'halt': False, 'reason': None, 'cost': 0.0}

--

## Layer 1: Tactical Layer

### Purpose

Pattern matching and simple rules. Suppress Layer 2 if high-confidence match found.

### ***Layer 1 Tactical Pattern Matching***

**Purpose: Match against learned patterns before calling expensive LLM**

**This layer can SUPPRESS Layer 2 (LLM) if confidence is high enough**

```
import re
```

```
class Layer1Tactical:
```

----

Layer 1 handles pattern-based decision-making:

- Pattern matching using regex (extremely fast)
- Rule engine evaluation (simple if-then logic)
- Cache lookup (avoid reprocessing same content)
- High-confidence pattern suppression (skip LLM if confident)

Cost: \$0.0001 per match (just cache lookup)

Speed: <10ms typical

Confidence threshold: 0.85 (configurable)

----

```
def __init__(self):
```

```
    # patterns: dict mapping pattern names to their configurations
```

```
    # Example: {'urgent': {'regex': r'urgent|asap', 'category': 'urgent', 'confidence': 0.95}}
```

```
    self.patterns = {}
```

```
    # confidence_threshold: minimum confidence needed to suppress Layer 2
```

```
    self.confidence_threshold = 0.85
```

```
# cache: dict mapping content hashes to previous results
self.cache = {}
```

```
def add_pattern(self, name, regex, category, confidence=0.92):
```

```
    """
```

```
    Register new pattern for matching
```

Args:

    name: string identifier for pattern (e.g., 'urgent')

    regex: compiled or raw regex pattern (e.g., r'urgent|asap')

    category: string category to assign if pattern matches

    confidence: float 0.0-1.0 indicating pattern accuracy

```
    """
```

```
    self.patterns[name] = {
```

        'regex': regex,

        'category': category,

        'confidence': confidence

```
}
```

```
def match_patterns(self, content):
```

```
    """
```

```
    Try to match content against all known patterns
```

Args:

    content: string text to classify

Returns:

    dict with:

        - matched: bool (True if pattern found)

        - pattern: str pattern name that matched

        - category: str assigned category

        - confidence: float confidence of match

        - method: str 'pattern' (for tracking)

        - cost: float 0.0 (no API calls)

Algorithm:

1. Convert content to lowercase (case-insensitive)

2. For each registered pattern:

- a. Check if pattern's regex matches

- b. If matches AND confidence >= threshold:
  - Return match immediately
- 3. If no pattern matches, return matched=False

```
#####
# Convert to lowercase for case-insensitive matching
content_lower = content.lower()

# Iterate through all registered patterns
for pattern_name, pattern_config in self.patterns.items():
    regex = pattern_config['regex']

    # Try to match pattern against content
    if re.search(regex, content_lower):
        confidence = pattern_config['confidence']

        # Check if confidence high enough to suppress LLM
        if confidence >= self.confidence_threshold:
            # Return successful match
            return {
                'matched': True,
                'pattern': pattern_name,
                'category': pattern_config['category'],
                'confidence': confidence,
                'method': 'pattern',
                'cost': 0.0
            }

# No patterns matched
return {'matched': False}
```

```
def check_cache(self, content_hash):
```

```
#####
Check if we've already processed this exact content
```

Args:

content\_hash: int hash of the content string

Returns:

dict with:

- matched: bool (True if found in cache)
- (plus all fields from original result if cached)

Purpose:

- Avoid re-processing identical content
- Reuse previous results without re-running LLM
- Reduce latency on duplicate requests

=====

```
if content_hash in self.cache:  
    cached = self.cache[content_hash]  
    # Mark as coming from cache for tracking  
    cached['method'] = 'cache'  
    cached['cost'] = 0.0  
    return {'matched': True, **cached}  
  
# Not in cache  
return {'matched': False}
```

def execute(self, request):

=====

Execute Layer 1 matching in optimal order

Args:

request: dict with 'content' key

Returns:

dict with:

- matched: bool (True if pattern/cache found result)
- category: str assigned category (if matched)
- confidence: float confidence (if matched)
- method: str 'cache' | 'pattern'
- cost: float 0.0 (Layer 1 never costs money)

Execution order:

1. Cache check first (fastest: hash lookup O(1))
2. Pattern matching second (fast: regex matching)
3. If neither works, return to Layer 2 (LLM call)

=====

content = request['content']

```

# Create hash of content for cache lookup
content_hash = hash(content)

# Strategy 1: Check cache first (fastest path)
cache_result = self.check_cache(content_hash)
if cache_result['matched']:
    # Cache hit! Return immediately
    return cache_result

# Strategy 2: Try pattern matching (second fastest path)
pattern_result = self.match_patterns(content)
if pattern_result['matched']:
    # Pattern matched! Cache this result for future
    self.cache[content_hash] = pattern_result
    return pattern_result

# Strategy 3: No cache or pattern match found
return {'matched': False, 'cost': 0.0}

```

---

**## Layer 2: Deliberative Layer**

**### Purpose**

Complex reasoning using LLM. Called only when Layer 1 cannot resolve.

**### Implementation**

### ***Layer 2 Deliberative LLM Reasoning***

**Purpose:** Call LLM for complex reasoning when simpler layers can't decide

class Layer2Deliberative:

=====

Layer 2 handles complex decision-making:

- Full LLM reasoning with context

- Novel decision-making (cases Layers 0-1 couldn't handle)
- Pattern learning (extract new patterns from LLM decisions)
- Context integration (use conversation history)

Cost: \$0.01-0.05 per call (expensive!)

Speed: 100-500ms typical

Confidence: 85-95% typical

This layer only reached when:

1. Layer 0 passed validation
  2. Layer 1 found no high-confidence pattern match
- .....

```
def __init__(self, llm_model='claude-3-5-sonnet-20241022', api_key=None):
```

.....

Initialize Layer 2 with LLM configuration

Args:

    llm\_model: string LLM model ID  
    api\_key: string API key (optional)

.....

```
    self.llm_model = llm_model
```

```
    self.api_key = api_key
```

```
    # cost_per_call: average cost of one LLM call
```

```
    self.cost_per_call = 0.01
```

```
    # learned_patterns: patterns discovered by this agent
```

```
    self.learned_patterns = {}
```

```
def call_llm(self, prompt, context=None):
```

.....

Call LLM for reasoning and categorization

Args:

    prompt: string question/task for LLM  
    context: optional dict with conversation history

Returns:

    string: LLM's response/decision

```
#####
# In production, this would call actual Anthropic API:
# client = anthropic.Anthropic(api_key=self.api_key)
# response = client.messages.create(...)

# For demo, use heuristic rules
prompt_lower = prompt.lower()

if 'urgent' in prompt_lower:
    return 'urgent'
elif 'invoice' in prompt_lower:
    return 'billing'
elif 'unsubscribe' in prompt_lower:
    return 'spam'
else:
    return 'other'
```

```
def learn_pattern(self, content, category, confidence):
#####
```

Extract and save new pattern from LLM decision

Args:

content: string original text that was categorized  
category: string category assigned by LLM  
confidence: float 0.0-1.0 confidence in decision

Purpose:

- Save high-confidence patterns to Layer 1
  - Next time we see similar content, Layer 1 handles it
  - Gradually shift expensive decisions to cheap layer
- ```
#####
```

```
# Only learn from high-confidence decisions
```

```
if confidence >= 0.90:
```

```
    # Extract key features from content
```

```
    words = content.lower().split()
```

```
    if words:
```

```
        # Use first word as pattern keyword
```

```
        key_word = words
```

```
# Save pattern
self.learned_patterns[category] = {
    'keyword': key_word,
    'confidence': confidence,
    'learned_at': __import__('datetime').datetime.now().isoformat()
}
```

def execute(self, request):

"""  
 Execute Layer 2 LLM reasoning

Args:

request: dict with 'content' key

Returns:

dict with:

- category: str assigned category
- method: str 'llm'
- confidence: float 0.85-0.95 typical
- cost: float \$0.01 (the expensive part!)
- layer: int 2 (for tracking)

"""

```
content = request['content']
```

# Create prompt for LLM

```
prompt = f"Categorize this text: {content}"
```

# Call LLM (expensive operation!)

```
category = self.call_llm(prompt)
```

# Learn pattern if high confidence

```
self.learn_pattern(content, category, 0.92)
```

# Return decision with full metadata

```
return {
```

- 'category': category,
- 'method': 'llm',
- 'confidence': 0.92,
- 'cost': self.cost\_per\_call,

```
'layer': 2  
}
```

---

```
## Integration: Three-Layer Pipeline
```

### **SPL Agent: Full Three-Layer Pipeline**

**Orchestrates all three layers in correct sequence**

```
class SPLAgent:
```

====

SPL Agent integrates all three layers:

- Layer 0: Validation (fast, free)
- Layer 1: Patterns (fast, cheap)
- Layer 2: LLM (slow, expensive)

Execution flow:

request → Layer 0 (halt?) → Layer 1 (match?) → Layer 2 (LLM) → result

Cost optimization:

- Most requests stop at Layer 0 or 1
- Only ~5-15% reach expensive Layer 2
- This achieves 10-50x cost reduction

====

```
def __init__(self):  
    # Initialize all three layers  
    self.layer0 = Layer0Reactive()  
    self.layer1 = Layer1Tactical()  
    self.layer2 = Layer2Deliberative()
```

```
# Pre-load common patterns (bootstrapping)  
self.layer1.add_pattern(
```

```
'urgent_keyword',
r'urgent|asap|immediately',
'urgent',
0.92
)

def process(self, request):
    """
    Process request through all three layers

    Args:
        request: dict with 'user_id' and 'content'

    Returns:
        dict with result and metadata
    
```

Execution stages:

=====

Stage 1: Layer 0 Validation

- If invalid → REJECT (\$0 cost)
- If valid → continue to Layer 1

Stage 2: Layer 1 Pattern Matching

- If pattern matches → RETURN (\$0 cost)
- If no match → continue to Layer 2

Stage 3: Layer 2 LLM

- Full reasoning → RETURN (\$0.01 cost)

====

```
# STAGE 1: Layer 0 Validation
layer0_result = self.layer0.execute(request)
if layer0_result['halt']:
    # Request rejected at validation stage
    return {
        'result': None,
        'reason': layer0_result['reason'],
        'cost': 0.0,
        'layer': 0
    }
```

```
    }

# STAGE 2: Layer 1 Pattern Matching
layer1_result = self.layer1.execute(request)
if layer1_result['matched']:
    # Pattern matched - suppress Layer 2!
    return {
        'result': layer1_result['category'],
        'confidence': layer1_result['confidence'],
        'method': layer1_result['method'],
        'cost': layer1_result['cost'],
        'layer': 1,
        'suppressed': True
    }
```

```
# STAGE 3: Layer 2 LLM Reasoning
layer2_result = self.layer2.execute(request)

return {
    'result': layer2_result['category'],
    'confidence': layer2_result['confidence'],
    'method': layer2_result['method'],
    'cost': layer2_result['cost'],
    'layer': 2
}
```

---

## Cost Optimization Strategies

### Strategy 1: Pattern Preloading

**\*Load high-confidence patterns at start**

```
def preload_patterns(agent):
    """Pre-populate agent with high-accuracy patterns"""
```

```
patterns = [
    ('urgent', r'urgent|asap|emergency', 'urgent', 0.95),
    ('billing', r'invoice|payment|bill|receipt', 'billing', 0.93),
    ('spam', r'unsubscribe|viagra|lottery|winner', 'spam', 0.98),
]
for name, regex, category, conf in patterns:
    agent.layer1.add_pattern(name, regex, category, conf)
```

### ### Strategy 2: Confidence Threshold Tuning

#### \*Adjust threshold based on domain

```
def set_confidence_threshold(agent, domain):
    """Configure confidence threshold for domain"""
    if domain == 'email_filtering':
        # High cost of false positives
        agent.layer1.confidence_threshold = 0.90 # Strict
    elif domain == 'content_moderation':
        # Balance between cost and accuracy
        agent.layer1.confidence_threshold = 0.85 # Moderate
    else:
        # Prioritize cost savings
        agent.layer1.confidence_threshold = 0.75 # Permissive
```

### ### Strategy 3: Timeout Protection

#### \*Prevent Layer 2 from running indefinitely

```

def execute_with_timeout(agent, request, timeout_ms=500):
    """Execute Layer 2 with timeout protection"""

    import time
    start = time.time()
    result = agent.layer2.execute(request)
    elapsed = (time.time() - start) * 1000

    if elapsed > timeout_ms:
        # Timeout occurred - fall back to Layer 1
        return agent.layer1.execute(request)

    return result

```

--

## Performance Benchmarks

### Single Request Processing

| Layer | Operation            | Time  | Cost   |
|-------|----------------------|-------|--------|
| 0     | Validation           | <1ms  | \$0    |
| 1     | Pattern match (hit)  | 2ms   | \$0    |
| 1     | Pattern match (miss) | 2ms   | \$0    |
| 2     | LLM call             | 200ms | \$0.01 |

### 1000 Email Batch

| Scenario        | Layer 0 | Layer 1 | Layer 2 | Total Cost | Avg/Email |
|-----------------|---------|---------|---------|------------|-----------|
| No optimization | 0       | 0       | 1000    | \$10.00    | \$0.01    |
| 50% patterns    | 0       | 500     | 500     | \$5.00     | \$0.005   |
| 80% patterns    | 10      | 800     | 190     | \$1.90     | \$0.0019  |
| 95% patterns    | 50      | 950     | 0       | \$0.00     | \$0.0000  |

--

```
## Monitoring & Observability
```

### \*Track performance metrics for optimization

```
class SPLMonitor:  
    """Collect metrics for performance monitoring"""
```

```
def __init__(self):  
    # Metrics dictionary  
    self.metrics = {  
        'layer0_halts': 0,  
        'layer1_hits': 0,  
        'layer2_calls': 0,  
        'total_cost': 0.0,  
        'total_requests': 0  
    }
```

```
def record_execution(self, result):  
    """Record metrics from one execution"""  
    self.metrics['total_requests'] += 1  
    self.metrics['total_cost'] += result['cost']  
  
    if result['layer'] == 0:  
        self.metrics['layer0_halts'] += 1  
    elif result['layer'] == 1:  
        self.metrics['layer1_hits'] += 1  
    elif result['layer'] == 2:  
        self.metrics['layer2_calls'] += 1
```

```
def report(self):
    """Generate performance report"""
    total = self.metrics['total_requests']
    return {
        'total_requests': total,
        'layer0_halt_rate': self.metrics['layer0_halts'] / total if total > 0 else 0,
        'layer1_hit_rate': self.metrics['layer1_hits'] / total if total > 0 else 0,
        'layer2_call_rate': self.metrics['layer2_calls'] / total if total > 0 else 0,
        'total_cost': self.metrics['total_cost'],
        'avg_cost_per_request': self.metrics['total_cost'] / total if total > 0 else 0
    }
```

---

## ## Deployment Checklist

- [ ] Layer 0 validation rules configured
- [ ] Layer 1 patterns loaded and tested
- [ ] Layer 2 LLM API credentials configured
- [ ] Rate limiting configured
- [ ] Monitoring/logging enabled
- [ ] Cost tracking enabled
- [ ] Timeout protection in place
- [ ] Error handling for all layers
- [ ] Fallback mechanisms configured
- [ ] Performance baselines established

---

\*Pamela Cuce | pamcuce@alum.vassar.edu | Version 3.1 (Corrected) | December 3, 2025\*

### ### Implementation

```
class Layer1Tactical:  
    def init(self):  
        self.patterns = {}  
        self.confidence_threshold = 0.85  
        self.cache = {}  
  
    def add_pattern(self, name, regex, category, confidence=0.92):  
        """Register a pattern"""  
        self.patterns[name] = {  
            'regex': regex,  
            'category': category,  
            'confidence': confidence  
        }  
  
    def match_patterns(self, content):  
        """Try to match against known patterns"""  
        content_lower = content.lower()  
  
        for pattern_name, pattern_config in self.patterns.items():  
            regex = pattern_config['regex']  
  
            if re.search(regex, content_lower):  
                confidence = pattern_config['confidence']
```

```

        if confidence >= self.confidence_threshold:
            return {
                'matched': True,
                'pattern': pattern_name,
                'category': pattern_config['category'],
                'confidence': confidence,
                'method': 'pattern',
                'cost': 0.0
            }

    return {'matched': False}

def check_cache(self, content_hash):
    """Check if result already cached"""
    if content_hash in self.cache:
        cached = self.cache[content_hash]
        cached['method'] = 'cache'
        cached['cost'] = 0.0
        return {'matched': True, **cached}

    return {'matched': False}

def execute(self, request):
    """Execute Layer 1 matching"""
    content = request['content']
    content_hash = hash(content)

    # Try cache first
    cache_result = self.check_cache(content_hash)
    if cache_result['matched']:
        return cache_result

    # Try pattern matching
    pattern_result = self.match_patterns(content)
    if pattern_result['matched']:
        # Cache this result
        self.cache[content_hash] = pattern_result
        return pattern_result

```

```
return {'matched': False, 'cost': 0.0}
```